

Craftax runs compared: achievements and what unlocks them

13 finished runs, 1.1M env steps each, size50m, one training seed per run. Scored with tools/death_eval.py on the same fixed evaluation worlds. Updated 2026-10-06.

Every number is the mean episode-final achievement count over the evaluation episodes, and each unlock rate is the share of episodes in which that achievement fired. Every run was scored on the same fixed evaluation worlds; re-scoring two old checkpoints with today's code reproduced them exactly.

The two-level agent sets the record. Manager v1.2 -- B2 plus a manager that picks a tech-tree goal and holds it until reached (or 32 steps), with one critic shared by both levels: a game-reward head judges both, a goal-reward head only the bottom actor -- reaches 9.73 achievements, against 8.53 for v1.1, 7.60 for B2 and 7.13 for v1. It is the first run to make the stone pickaxe regularly (23%), and it mines coal (10%), makes both wood tools far more often (wood sword 57%, B2 20%) and places a table in every episode. v1.1 (8.53), which pays both levels when a goal is reached, places the furnace and stone in 67% of episodes and makes a stone sword (23%) and arrows (10%); v1, whose manager never learned to choose, stays below B2.

Experiment B before it: masking impossible actions (B2) 7.60 against 5.87 for the same model without it; validity flags as an input (B1) 6.90. The non-cheating model (Honest map) matches the privileged control (5.87 vs 5.93).

Every map run except the hand-written-rewards one unlocks WAKE_UP (sleep, then wake) in 80-98% of episodes; vanilla never does.


Survival is still unsolved and v1.2 dies earliest (221 steps): it does more, and half its deaths are mobs or lava. Thirst is the first meter to run out in about half of every manager run's deaths. In normalized return v1.2 reaches 3.9% of the 226-point maximum (v1.1 3.4%, B2 3.0%), against 7.09% for the best published 1M-step agent (ITC). The masking and the manager's goals use game knowledge at test time, which the published agents do not.

run	eps	ach.	w/o wake	life	table	pickaxe	sword	stone	drink
Vanilla DreamerV3	25	4.44	4.44	262	64%	24%	0%	12%	80%
Map model (bug)	60	4.35	3.37	294	35%	0%	0%	0%	43%
Map + 5 hand rewards (bug)	50	2.30	2.30	263	0%	0%	0%	0%	44%
Map + potential (bug)	30	3.87	2.94	277	13%	0%	0%	0%	37%
Map model (fixed)	30	5.03	4.20	264	63%	7%	0%	0%	70%
Map + potential (fixed)	30	5.80	4.97	278	90%	30%	10%	20%	60%
Control: + recipe & death fixes	30	5.93	5.03	268	83%	0%	60%	0%	67%
Honest map (observation-only)	30	5.87	4.90	243	87%	20%	20%	7%	83%
B1: validity flags as input	30	6.90	6.10	311	87%	50%	30%	33%	80%
B2: impossible actions masked	30	7.60	6.80	246	87%	53%	20%	43%	70%
Manager v1: B2 + goal manager	30	7.13	6.20	254	73%	57%	23%	47%	63%
Manager v1.1: + reach bonuses	30	8.53	7.66	259	93%	73%	33%	73%	57%
Manager v1.2: shared critic	30	9.73	8.90	221	100%	80%	57%	77%	57%


eps = evaluation episodes. ach. = mean episode-final achievements; w/o wake = the same minus WAKE_UP. life = mean lifespan in steps. Rates are % of episodes. 'bug' = imag_shift True, which corrupted the policy gradient (fixed 2026-09-20). map_survive was re-evaluated on 2026-09-28 and reproduced its original 2.30. One training seed per run, so the run-to-run spread is unknown. 'Map + potential (fixed)' is the old 5.80 record; 'Control' adds the recipe and death fixes with the privileged map; 'Honest map' has the same fixes with observation-only targets. 'B1' and 'B2' are experiment B on the honest-map configuration: validity flags as an encoder input, and impossible actions masked in training. 'mgr v1' and 'mgr v1.1' are B2 plus the two-level agent (docs/design-manager.md): v1 pays the manager the game reward only; v1.1 also pays both levels when the goal is reached and masks goals that already hold; v1.2 shares one two-headed critic between both levels and holds each goal until reached or 32 steps.

Sleep inflates the map runs; the manager is real progress

Mean episode-final achievements per run, split into WAKE_UP and the rest; unlock rate of the early tech-tree rungs



Unlock rate (% of episodes), early tech tree in dependency order




Every achievement any run unlocked

% of evaluation episodes in which each achievement fired. The other 47 of 67 were never unlocked by any run.

	vanilla	map, bug	map+5 terms	map+ $\emptyset$, bug	map	map+ $\emptyset$	map+ $\emptyset$ +fixes	honest map	B1 told	B2 masked	mgr v1	mgr v1.1	mgr v1.2
collect wood	76	72	52	63	83	100	97	93	100	93	93	97	100
place table	64	35	.	13	63	90	83	87	87	87	73	93	100
eat cow	4	12	22	13	10	3	17	3	23	13	3	13	7
collect sapling	92	92	98	90	93	93	93	90	93	97	93	90	100
collect drink	80	43	44	37	70	60	67	83	80	70	63	57	57
make wood pickaxe	24	.	.	.	7	30	.	20	50	53	57	73	80
make wood sword	.	.	.	.	.	10	60	20	30	20	23	33	57
place plant	92	77	.	73	93	87	83	87	83	93	70	67	100
defeat zombie	.	5	12	3	.	.	3	.	7	3	.	3	10
collect stone	12	.	.	.	.	20	.	7	33	43	47	73	77
place stone	.	.	.	.	.	3	.	.	.	40	43	67	73
defeat skeleton	.	2	2	.	.	.	.	.	.	.	.	.	.
make stone pickaxe	.	.	.	.	.	.	.	.	.	.	3	.	23
make stone sword	.	.	.	.	.	.	.	.	.	.	10	23	37
wake up	.	98	.	93	83	83	90	97	80	80	93	87	83
place furnace	.	.	.	.	.	.	.	.	23	40	40	67	60
collect coal	.	.	.	.	.	.	.	.	.	20	.	.	10
make arrow	.	.	.	.	.	.	.	.	.	.	.	10	.
make torch	.	.	.	.	.	.	.	.	.	3	.	.	.
place torch	.	.	.	.	.	.	.	.	.	3	.	.	.

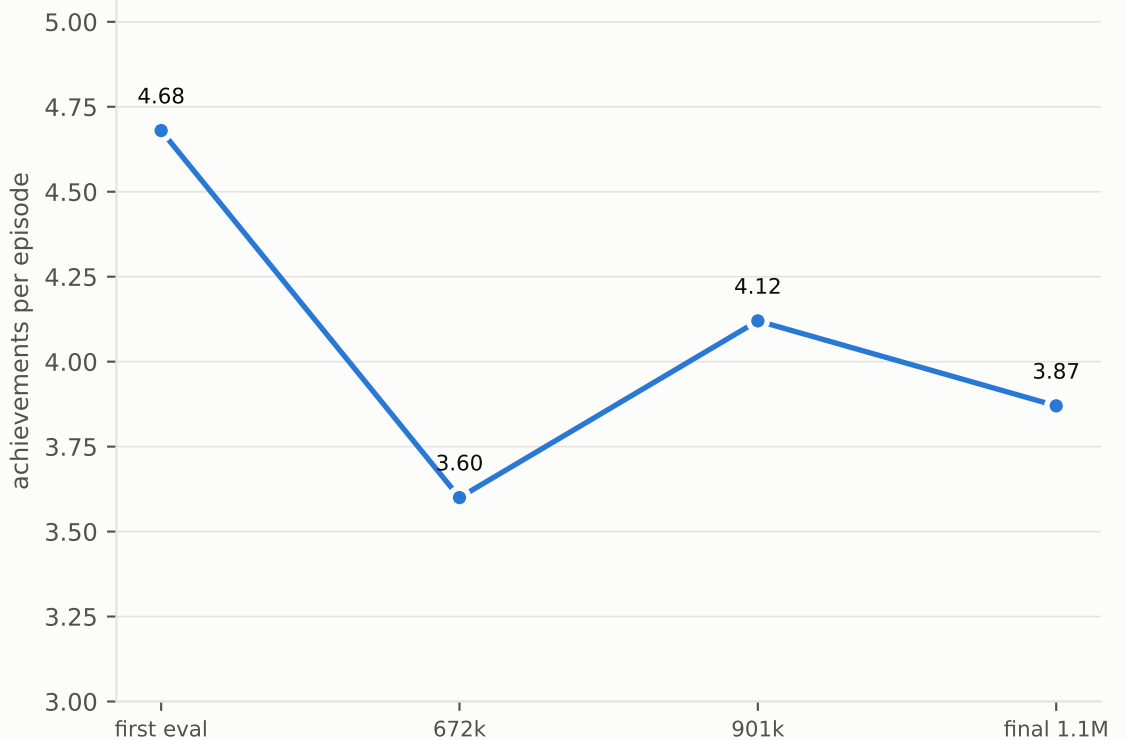
Survival never moved

Cause of death (share of evaluation episodes) and lifespan; and how much one run moves between its own checkpoints





One run, four checkpoints: Map + potential (bug)

The same training run, scored at four points (25-30 episodes each), moves by 1.08 achievements. That is the scale of noise any comparison between single runs has to beat.



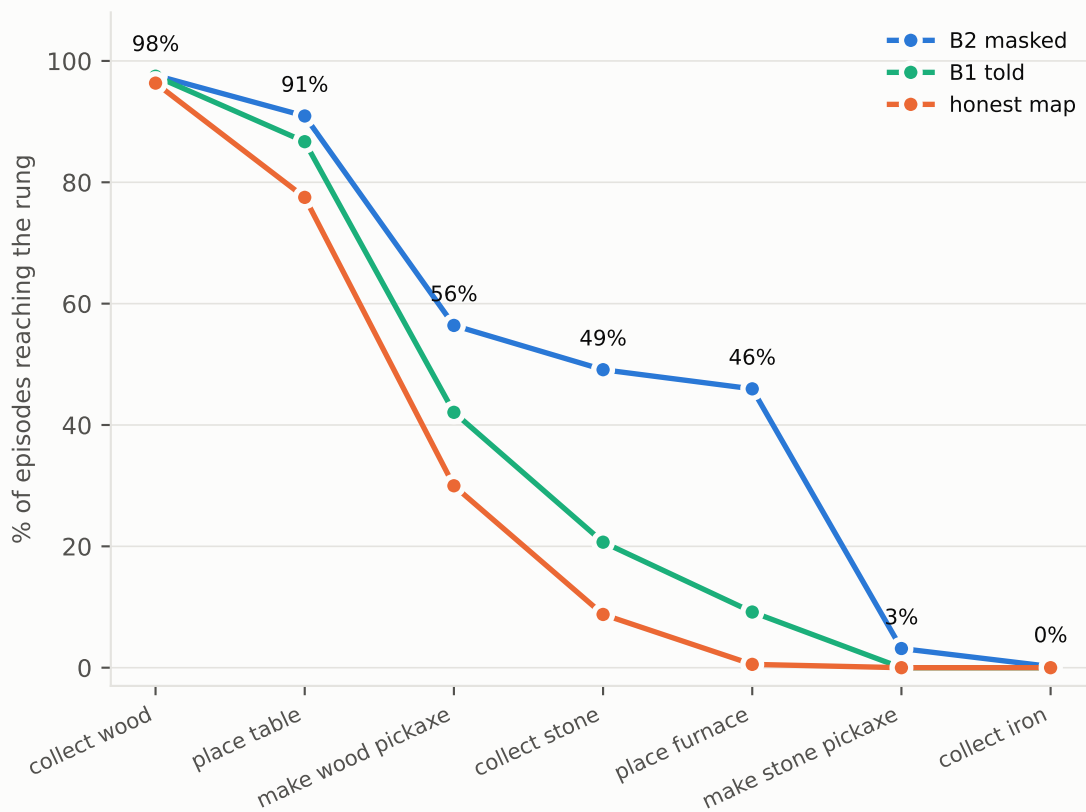
Training curves: every run

Rebuilt from each run's replay; 50k-step bins; value = mean of the last 100k steps; one seed per run.



Where B2 falls short

Last 150k training steps of each run (~500 episodes), decoded from replay by tools/diagnose_run.py



1. The stone pickaxe is the wall. B2 reaches the furnace in 46% of episodes but makes the stone pickaxe in 3%. Of the 249 episodes that collected stone, 73% held wood and stone at the same time, but only 22% ever stood next to a table while holding both, and 6% made the pickaxe. Two gaps: it rarely brings the ingredients back to a table (logistics), and when it does the chance lasts one step -- the single log usually goes into the wood sword or it walks off (measured in docs/b2-casebook.pdf).
2. Table to wood pickaxe leaks. 38% of the episodes that place a table never make the wood pickaxe. At a table holding wood it crafts every time, but every one of these episodes placed the table with its last two logs, and later logs go into another table rather than back to the first (docs/b2-casebook.pdf).
3. Survival. Training-window lifespan 285 (median 275) against 261-284 for a random policy. Thirst is still the first meter to run out in 45% of deaths and hunger in 24%: Crafttax pays for the first drink only, so topping up is never rewarded. A hostile mob is adjacent at 61% of deaths; it dies in lava in 4% (the others: under 1%), the price of exploring further. Deaths in its sleep are rare (5%).
4. Wasted presses. 43% of presses still do nothing (experiment A): DO with nothing in front (26%), NOOP and walking into walls. Masking only removes the special actions; these basic ones are never masked.

What the results say, and what would settle the rest

0. The manager pays off once it learns to choose -- most with a shared critic

v1 paid the manager the game reward only; most goals cannot be reached in 8 steps, so the actor ignored them, the manager stayed uniform, and it scored below B2 (7.13). v1.1 pays both levels when a goal is reached; its manager learned to choose and it scored 8.53. v1.2 replaces the manager's own critic with one two-headed critic shared with the bottom actor, normalises the game and goal rewards separately, and holds each goal until reached: 9.73, and the first regular stone pickaxe (23%). Its manager favours goals it cannot complete yet (iron pickaxe, iron sword, diamond), which pay the bottom actor for gathering every ingredient on the way -- plausibly why it stocks up instead of spending its wood at once. All three trailed B2 early and overtook it between 350k and 550k steps. One seed per run.

1. Masking impossible actions is the largest gain so far

B2 7.60 vs honest map 5.87, same model and reward. Experiment A on its final checkpoint shows why: buttons that sat at exactly 0% in the states where they would work now get pressed there (furnace 3.2% over 599 states, placing stone 2.4%, placing a table 75%, the wood pickaxe 100%), and every no-effect press is now a basic action (DO, NOOP, a blocked move). Once a useless press can no longer teach the agent that a button is bad, it presses that button when it works.

2. Telling helps, masking helps more

B1 (flags as input) reaches 6.90: both wood crafts survive in one run for the first time and the furnace comes off zero, but placing stone stays at 0.0% across 452 chances and about 6% of presses still go to impossible buttons. The information helps; removing the suppressing gradient helps more. Masking changes the action space relative to published Craftax baselines, so it is reported as a separate arm.

3. The next wall is sequencing, not suppression

In B2 the stone pickaxe was craftable in only 5 of 5,167 evaluation steps: the agent mines stone but rarely returns to a table holding both wood and stone. That gates iron and every 3-5-point achievement after it. It is a planning problem, so the next experiment should target multi-step sequences (experiment D, return-then-explore, or a sub-goal signal), not the action distribution.

4. Survival is still unsolved

Lifespans stay at or below the random band, except B1 (311, mean pulled up by a few long episodes; median 265). The better runs die earlier: B2 246, v1.2 221 steps, with half of v1.2's deaths to mobs or lava -- the more it does, the more it is exposed -- and thirst first to run out in most of the rest. Manager v1.3 targets this with survival goals and a survival potential.

5. What to run next

Training now: manager v1.3 (v1.2 plus survival goals and a survival potential, aimed at the thirst deaths). Queued: B3 (B2 with the action mask learned instead of transcribed from the rules, so no game rules at test time). Then the best manager version with B3's learned mask, and second seeds of the leaders, since every comparison here is single-seed.